

Stakeholder Engagement

Stakeholder map and engagement plan for the harmful-intent detection project

Project	Detecting harmful/malicious user intent via internal LLM representations
Team	Komadugu Team
Document status	Engagement plan — proposed methods and cadence for this prototype stage

1. Purpose

This document maps the people and groups affected by, or able to influence, the project, and lays out how the team plans to engage them as the work moves from a research prototype toward anything closer to real-world use. At this stage the project is a challenge-scale prototype (Gemma-2-2B probe trained on a public benchmark dataset); no live users or partner organizations have been engaged yet, so this document is written as an honest **plan** rather than a record of consultations already completed.

2. Stakeholder Map

Stakeholder	Stake / Interest	Influence	Engagement Priority
Everyday AI users	Want safe, trustworthy responses without being wrongly blocked.	Low (individually) / High (collectively, via feedback and adoption)	High
Organizations deploying LLM apps (support, education, healthcare, finance)	Need a moderation layer that is accurate, fast, and low-cost to integrate.	High — decide whether/how to adopt	High
AI developers & safety researchers	Want interpretable, reproducible methods they can validate and build on.	High — technical review, replication	High
Online communities / moderators	Experience the downstream effects of over- or under-flagging.	Medium	Medium
Vulnerable / marginalized groups	Disproportionately exposed to harm if detection fails, or to unfair flagging if biased.	Low direct influence, high impact	High
Dataset originators (Jigsaw / Conversation AI, original commenters)	Interest in how their released data / historical comments are reused.	Low (platform is defunct)	Low–Medium
Course instructors / mentors / evaluators	Assess technical soundness and responsible-AI process.	High for this deliverable	High

3. Engagement Methods by Stage

Current stage (prototype / challenge submission)

- Grounding the problem statement in documented, publicly reported harms of LLM misuse rather than assumption (desk research).
- Using a dataset (Civil Comments) that itself encodes thousands of real annotator judgments about what reads as toxic, as an indirect proxy for public perspective.
- Internal review against the team's values-led problem statement at each major design decision (dataset choice, threshold, evaluation metric).
- Documenting known limitations transparently in the Data Card and Impact Statement Card so reviewers/mentors can challenge assumptions.

Next stage (pilot / limited deployment)

- Structured feedback sessions with a pilot deploying organization (e.g., a support or education tool) to review false positive/negative examples together.
- Lightweight user-facing feedback mechanism (e.g., a "this was flagged incorrectly" control) surfaced through the Streamlit interface or an integrating app.
- Targeted review with representatives of groups most affected by identity-term bias in toxicity classifiers, informed by the Data Card's documented limitation.
- Sharing evaluation results (including subgroup breakdowns, once available) with AI safety researchers for external critique.

Ongoing (post-adoption)

- Periodic re-evaluation as language and threats evolve, with a defined owner and cadence (e.g., quarterly).
- A documented escalation path for moderators/users who believe a decision was wrong, feeding into re-audits.
- Publishing updated cards (Data Card, Impact Statement Card) whenever the model, data, or threshold changes materially.

4. Feedback Loop

Feedback collected from any stage above is intended to flow back into three concrete levers already built into the pipeline: (1) the toxicity **threshold** and **class-weighting** used in training (`src/config.py`, `src/train.py`), (2) the **evaluation protocol** — adding subgroup-disaggregated metrics alongside overall accuracy, and (3) the **documentation** itself (this plan and the two accompanying cards), which should be revised whenever a stakeholder surfaces a limitation not yet captured here.

5. Open Gaps & Next Steps

- No real end users or partner organizations have been consulted yet — the highest-priority next step before any pilot.
- Subgroup fairness evaluation (by identity/topic) has not yet been run on this project's specific validation split.
- No formal, named point of contact for external feedback has been established yet.
- Multilingual stakeholders are currently out of scope given the English-only training data.